

1. Accomplishments to date

- Determined the data imbalance is 99.8% for normal transactions and 0.2% fraudulent transactions. Also created a bar plot to visually show just how large the imbalance is.
- To reduce training and testing time I took a sample of 10% from the original features and classes to train my models with.
- Then I performed an initial train, test, split with the testing size being 20%. Next, I repeated the train, test, split for validation.
- Created a KNeighborsClassifier and used GridSearchCV with 5 fold cross validation. This resulted in an accuracy score of 100%.
- Created a DecisionTreeClassifier and used GridSearchCV with 5 fold cross validation. This resulted in an accuracy score of 100%.
- Created a GaussianNB and used GridSearchCV with 5 fold cross validation. This resulted in an accuracy score of 100%.
- Created a LinearRegression and used GridSearchCV with 5 fold cross validation. This resulted in an accuracy score of 100%.
- Created a LogisticRegression and used GridSearchCV with 5 fold cross validation. This resulted in an accuracy score of 100%.
- Created a LinearSVC and used GridSearchCV with 5 fold cross validation. This resulted in an accuracy score of 100%.
- Created a Keras neural network using all 30 features in the first layer 1 hidden layer with 50 nodes and an output layer with 1 node.

2. Remaining steps

- Need to check models for over fitting
- Find a better metric to use for imbalanced data
- Create better visuals to illustrate models and data
- Fix neural network model and tensorboard

Code:

```
'''  
Alex Lux
```

The challenge is to recognize fraudulent credit card transactions so that the customers of credit card companies are not charged for items that they did not purchase.

Main challenges involved in credit card fraud detection are:

1. Enormous Data is processed every day and the model build must be fast enough to respond to the scam in time.
2. Imbalanced Data i.e most of the transactions (99.8%) are not fraudulent which makes it really hard for detecting the fraudulent ones
3. Data availability as the data is mostly private.
4. Misclassified Data can be another major issue, as not every fraudulent transaction is caught and reported.
5. Adaptive techniques used against the model by the scammers.

How to tackle these challenges?

1. The model used must be simple and fast enough to detect the anomaly and classify it as a fraudulent transaction as quickly as possible.
2. Imbalance can be dealt with by properly using some methods which we will talk about in the next paragraph
3. For protecting the privacy of the user the dimensionality of the data can be reduced.
4. A more trustworthy source must be taken which double-check the data, at least for training the model.
5. We can make the model simple and interpretable so that when the scammer adapts to it with just some tweaks we can have a new model up and running to deploy.

```
'''
```

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
LABELS = ["Normal", "Fraud"]
```

```
df = pd.read_csv("creditcard.csv")
```

```
print("=====")  
print("=====")  
print("===== DATA FRAME")  
print("=====")  
print("=====")  
print("=====")
```

```

print(df.head())
print("=====
=====")
print("===== DATA DESCRIPTION
=====")
print("=====
=====")
print(df.describe())
print()
print("===== NULL VALUES
=====")
print(df.isnull().values.any())

count_classes = pd.value_counts(df['Class'], sort = True)
print(print("===== CLASS COUNT (99.8% Normal, 0.2%
fraud) ====="))
print(count_classes)
count_classes.plot(kind = 'bar', rot=0)
plt.title("Transaction Class Distribution")
plt.xticks(range(2), LABELS)
plt.xlabel("Class")
plt.ylabel("Frequency")
plt.show()

X = df.drop(['Class'], axis = 1)

y = df['Class']

X_data = X.values
y_data = y.values

X_data_sample = X.sample(frac=0.1, random_state=123)
y_data_sample = y.sample(frac=0.1, random_state=123)

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X_data_sample, y_data_sample,
test_size=0.2, random_state=123)

# print("X train:\n", X_train)
# print("y train:\n", y_train)

import warnings
warnings.filterwarnings('ignore', category=UserWarning)
import sys, os

if not sys.warnoptions:

```

```
warnings.simplefilter("ignore")
os.environ["PYTHONWARNINGS"] = "ignore" # Also affect subprocesses

X_train_new, X_valid, y_train_new, y_valid = train_test_split(X_train, y_train,
test_size=0.2, random_state=123)

from sklearn.model_selection import GridSearchCV
from sklearn.metrics import accuracy_score, classification_report

from sklearn.neighbors import KNeighborsClassifier
parameters_knn = {'n_neighbors':[1, 3, 5, 7, 9, 11, 13, 15], 'metric': ["manhattan",
"chebyshev", "hamming"], 'weights': ["uniform", "distance"]}
knn = KNeighborsClassifier()
knn_clf = GridSearchCV(estimator=knn, param_grid=parameters_knn, cv=5, n_jobs=-1,
verbose=1)
knn_clf.fit(X_train_new, y_train_new)
print(f"BEST PARAMETERS FOR KNN CLASSIFIER: {knn_clf.best_params_}")
knn_best = knn_clf.best_estimator_
knn_predictions = knn_best.predict(X_valid)
print(classification_report(y_valid, knn_predictions))

from sklearn.tree import DecisionTreeClassifier
parameters_dt = {"criterion": ["gini", "entropy"], "splitter": ["best", "random"]}
dt = DecisionTreeClassifier()
dt_clf = GridSearchCV(estimator=dt, param_grid=parameters_dt, cv=5, n_jobs=-1,
verbose=1)
dt_clf.fit(X_train_new, y_train_new)
print(f"BEST PARAMETERS FOR DECISION TREE CLASSIFIER: {dt_clf.best_params_}")
dt_best = dt_clf.best_estimator_
dt_predictions = dt_best.predict(X_valid)
print(classification_report(y_valid, dt_predictions))

from sklearn.naive_bayes import GaussianNB
nb = GaussianNB()
params_NB = {'var_smoothing': np.logspace(0,-9, num=100)}
gs_NB_clf = GridSearchCV(estimator=nb, param_grid=params_NB, cv=5, n_jobs=-1,
verbose=1)
gs_NB_clf.fit(X_train_new, y_train_new)
print(f"BEST PARAMETERS FOR GAUSSIAN NAIVE BAYES CLASSIFIER:
{gs_NB_clf.best_params_}")
nb_best = gs_NB_clf.best_estimator_
nb_predictions = nb_best.predict(X_valid)
print(classification_report(y_valid, nb_predictions))

from sklearn.linear_model import LinearRegression
lr = LinearRegression()
```

```

parameters_lr = {'fit_intercept':[True,False], 'copy_X':[True, False]}
lr_clf = GridSearchCV(estimator=lr, param_grid=parameters_lr, cv=5, n_jobs=-1,
verbose=1)
lr_clf.fit(X_train_new, y_train_new)
print(f"BEST PARAMETERS FOR LINEAR REGRESSION CLASSIFIER: {lr_clf.best_params_}")
lr_best = lr_clf.best_estimator_
lr_predictions = lr_best.predict(X_valid)
print(classification_report(y_valid, lr_predictions.round()))

```

```

from sklearn.linear_model import LogisticRegression
logistic_r = LogisticRegression()
parameters_logistic_r = {'penalty':['none', 'l2', 'l1', 'elasticnet'], 'C':[0.01, 0.1,
0.5, 1], 'solver': ['newton-cg', 'lbfgs', 'liblinear', 'sag', 'saga']}
logistic_r_clf = GridSearchCV(estimator=logistic_r, param_grid=parameters_logistic_r,
cv=5, n_jobs=-1, verbose=1)
logistic_r_clf.fit(X_train_new, y_train_new)
print(f"BEST PARAMETERS FOR LINEAR REGRESSION CLASSIFIER:
{logistic_r_clf.best_params_}")
logistic_r_best = logistic_r_clf.best_estimator_
logistic_r_predictions = logistic_r_best.predict(X_valid)
print(classification_report(y_valid, logistic_r_predictions))

```

```

from sklearn.svm import LinearSVC
svc = LinearSVC()
parameters_svc = {'penalty': ['l1', 'l2', 'elasticnet', 'none'], 'loss': ['hinge',
'squared_hinge'], 'C': [1, 10, 0.1]}
svc_clf = GridSearchCV(estimator=svc, param_grid=parameters_svc, cv=3, n_jobs=-1,
verbose=1)
svc_clf.fit(X_train_new, y_train_new)
print(f"BEST PARAMETERS FOR LinearSVC CLASSIFIER: {svc_clf.best_params_}")
svc_best = svc_clf.best_estimator_
svc_predictions = svc_best.predict(X_valid)
print(classification_report(y_valid, svc_predictions))

```

```

from tensorflow import keras
import tensorflow as tf
from keras.layers import Dense
from keras.models import Sequential

```

```

# Number of features (first layer inputs)
n_inputs = 30

```

```

nn_model = Sequential()
# define first hidden layer and visible layer
nn_model.add(Dense(50, input_dim=n_inputs, activation='relu',
kernel_initializer='he_uniform'))

```

```

# define output layer
nn_model.add(Dense(1, activation='sigmoid'))
# define loss and optimizer
nn_model.compile(loss='binary_crossentropy', optimizer='adam')

import datetime

log_dir = "logs/" + datetime.datetime.now().strftime("%Y-%m-%d-%H_%M_%S")
filepath = 'nn_model.hdf5'

from keras.callbacks import ModelCheckpoint

checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_accuracy', verbose=3,
save_best_only=True, mode='max')
tensorboard_callbacks = tf.keras.callbacks.TensorBoard(log_dir=log_dir,
histogram_freq=1)

nn_model.fit(X_train_new, y_train_new, epochs=10, callbacks=[checkpoint,
tensorboard_callbacks])
eval = nn_model.evaluate(X_valid, y_valid)
print(f"EVALUATION: {eval}")

nn_predictions = nn_model.predict(X_valid)

from sklearn.metrics import roc_auc_score
print(roc_auc_score(y_valid, nn_predictions))

nn_predictions_flat = nn_predictions.flatten()
y_pred = np.where(nn_predictions_flat > 0.5, 1, 0)

print(accuracy_score(y_valid, y_pred))
print(classification_report(y_valid, y_pred))

from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error

r2Score = r2_score(y_valid, y_pred)
maeScore = mean_absolute_error(y_valid, y_pred)
mseScore = mean_squared_error(y_valid, y_pred)
print(f"R2: {r2Score}, MAE: {maeScore}, MSE: {mseScore}")

```

Terminal Output:

```
PS C:\Users\Alex Lux\Desktop\DATA MINING\Project> python3 .\kaggleCCF.py
```

```
=====
===== DATA FRAME =====
=====
Time    V1    V2    V3    V4    V5    V6    V7    V8    V9    V10 ...  V20    V21    V22    V23    V24    V25    V26
V27    V28 Amount Class
0  0.0 -1.359807 -0.072781 2.536347 1.378155 -0.338321 0.462388 0.239599 0.098698 0.363787 0.090794 ... 0.251412 -0.018307 0.277838 -
0.110474 0.066928 0.128539 -0.189115 0.133558 -0.021053 149.62
0
1  0.0 1.191857 0.266151 0.166480 0.448154 0.060018 -0.082361 -0.078803 0.085102 -0.255425 -0.166974 ... -0.069083 -0.225775 -0.638672
0.101288 -0.339846 0.167170 0.125895 -0.008983 0.014724 2.69
0
2  1.0 -1.358354 -1.340163 1.773209 0.379780 -0.503198 1.800499 0.791461 0.247676 -1.514654 0.207643 ... 0.524980 0.247998 0.771679
0.909412 -0.689281 -0.327642 -0.139097 -0.055353 -0.059752 378.66
0
3  1.0 -0.966272 -0.185226 1.792993 -0.863291 -0.010309 1.247203 0.237609 0.377436 -1.387024 -0.054952 ... -0.208038 -0.108300 0.005274 -
0.190321 -1.175575 0.647376 -0.221929 0.062723 0.061458 123.50
0
4  2.0 -1.158233 0.877737 1.548718 0.403034 -0.407193 0.095921 0.592941 -0.270533 0.817739 0.753074 ... 0.408542 -0.009431 0.798278 -
0.137458 0.141267 -0.206010 0.502292 0.219422 0.215153 69.99
0
```

```
[5 rows x 31 columns]
```

```
=====
===== DATA DESCRIPTION =====
=====
Time    V1    V2    V3    V4    V5    V6 ...  V24    V25    V26    V27    V28    Amount
Class
count 284807.000000 2.848070e+05 2.848070e+05 2.848070e+05 2.848070e+05 2.848070e+05 2.848070e+05 2.848070e+05 ... 2.848070e+05 2.848070e+05
2.848070e+05 2.848070e+05 2.848070e+05 284807.000000 284807.000000
mean  94813.859575 1.168375e-15 3.416908e-16 -1.379537e-15 2.074095e-15 9.604066e-16 1.487313e-15 ... 4.473266e-15 5.340915e-16
1.683437e-15 -3.660091e-16 -1.227390e-16 88.349619 0.001727
std   47488.145955 1.958696e+00 1.651309e+00 1.516255e+00 1.415869e+00 1.380247e+00 1.332271e+00 ... 6.056471e-01 5.212781e-01
4.822270e-01 4.036325e-01 3.300833e-01 250.120109 0.041527
min   0.000000 -5.640751e+01 -7.271573e+01 -4.832559e+01 -5.683171e+00 -1.137433e+02 -2.616051e+01 ... -2.836627e+00 -1.029540e+01 -
2.604551e+00 -2.256568e+01 -1.543008e+01 0.000000 0.000000
25%   54201.500000 -9.203734e-01 -5.985499e-01 -8.903648e-01 -8.486401e-01 -6.915971e-01 -7.682956e-01 ... -3.545861e-01 -3.171451e-01 -
3.269839e-01 -7.083953e-02 -5.295979e-02 5.600000 0.000000
50%   84692.000000 1.810880e-02 6.548556e-02 1.798463e-01 -1.984653e-02 -5.433583e-02 -2.741871e-01 ... 4.097606e-02 1.659350e-02 -
5.213911e-02 1.342146e-03 1.124383e-02 22.000000 0.000000
75%   139320.500000 1.315642e+00 8.037239e-01 1.027196e+00 7.433413e-01 6.119264e-01 3.985649e-01 ... 4.395266e-01 3.507156e-01
2.409522e-01 9.104512e-02 7.827995e-02 77.165000 0.000000
max   172792.000000 2.454930e+00 2.205773e+01 9.382558e+00 1.687534e+01 3.480167e+01 7.330163e+01 ... 4.584549e+00 7.519589e+00
3.517346e+00 3.161220e+01 3.384781e+01 25691.160000 1.000000
```

```
[8 rows x 31 columns]
```

```
===== NULL VALUES =====
```

```
False
```

```
===== CLASS COUNT (99.8% Normal, 0.2% fraud) =====
```

```
None
```

```
0 284315
```

```
1 492
```

```
Name: Class, dtype: int64
```

```
Fitting 5 folds for each of 48 candidates, totalling 240 fits
```

```
BEST PARAMETERS FOR KNN CLASSIFIER: {'metric': 'manhattan', 'n_neighbors': 3, 'weights': 'uniform'}
```

```
precision recall f1-score support
```

```
0 1.00 1.00 1.00 4549
```

```
1 1.00 0.12 0.22 8
```

```
accuracy 1.00 4557
```

```
macro avg 1.00 0.56 0.61 4557
```

```
weighted avg 1.00 1.00 1.00 4557
```

```
Fitting 5 folds for each of 4 candidates, totalling 20 fits
```

```
BEST PARAMETERS FOR DECISION TREE CLASSIFIER: {'criterion': 'gini', 'splitter': 'random'}
```

```
precision recall f1-score support
```

```
0 1.00 1.00 1.00 4549
```

```
1 0.71 0.62 0.67 8
```

accuracy			1.00	4557
macro avg	0.86	0.81	0.83	4557
weighted avg	1.00	1.00	1.00	4557

Fitting 5 folds for each of 100 candidates, totalling 500 fits

BEST PARAMETERS FOR GAUSSIAN NAIVE BAYES CLASSIFIER: {'var_smoothing': 1.0}

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.00	1.00	1.00	4549
1	0.00	0.00	0.00	8

accuracy			1.00	4557
macro avg	0.50	0.50	0.50	4557
weighted avg	1.00	1.00	1.00	4557

Fitting 5 folds for each of 4 candidates, totalling 20 fits

BEST PARAMETERS FOR LINEAR REGRESSION CLASSIFIER: {'copy_X': True, 'fit_intercept': False}

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.00	1.00	1.00	4549
1	1.00	0.62	0.77	8

accuracy			1.00	4557
macro avg	1.00	0.81	0.88	4557
weighted avg	1.00	1.00	1.00	4557

Fitting 5 folds for each of 80 candidates, totalling 400 fits

BEST PARAMETERS FOR LINEAR REGRESSION CLASSIFIER: {'C': 0.01, 'penalty': 'none', 'solver': 'newton-cg'}

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.00	1.00	1.00	4549
1	0.71	0.62	0.67	8

accuracy			1.00	4557
macro avg	0.86	0.81	0.83	4557
weighted avg	1.00	1.00	1.00	4557

Fitting 3 folds for each of 24 candidates, totalling 72 fits

BEST PARAMETERS FOR LINEAR REGRESSION CLASSIFIER: {'C': 10, 'loss': 'hinge', 'penalty': 'l2'}

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.00	1.00	1.00	4549
1	0.00	0.00	0.00	8

accuracy			1.00	4557
macro avg	0.50	0.50	0.50	4557
weighted avg	1.00	1.00	1.00	4557

2022-05-06 17:21:33.656212: W tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic library 'cudart64_110.dll'; dlerror: cudart64_110.dll not found

2022-05-06 17:21:33.656356: I tensorflow/stream_executor/cuda/cudart_stub.cc:29] Ignore above cudart dlerror if you do not have a GPU set up on your machine.

2022-05-06 17:21:36.184215: W tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic library 'cudart64_110.dll'; dlerror: cudart64_110.dll not found

2022-05-06 17:21:36.184505: W tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic library 'cublas64_11.dll'; dlerror: cublas64_11.dll not found

2022-05-06 17:21:36.184840: W tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic library 'cublasLt64_11.dll'; dlerror: cublasLt64_11.dll not found

2022-05-06 17:21:36.185085: W tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic library 'cufft64_10.dll'; dlerror: cufft64_10.dll not found

2022-05-06 17:21:36.185425: W tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic library 'curand64_10.dll'; dlerror: curand64_10.dll not found

2022-05-06 17:21:36.185747: W tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic library 'cusolver64_11.dll'; dlerror: cusolver64_11.dll not found

2022-05-06 17:21:36.186021: W tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic library 'cusparse64_11.dll'; dlerror: cusparse64_11.dll not found

2022-05-06 17:21:36.186369: W tensorflow/stream_executor/platform/default/dso_loader.cc:64] Could not load dynamic library 'cudnn64_8.dll'; dlerror: cudnn64_8.dll not found

2022-05-06 17:21:36.186460: W tensorflow/core/common_runtime/gpu/gpu_device.cc:1850] Cannot dlopen some GPU libraries. Please make sure the missing libraries mentioned above are installed properly if you would like to use GPU. Follow the guide at <https://www.tensorflow.org/install/gpu> for how to download and setup the required libraries for your platform.

Skipping registering GPU devices...


```
2022-05-06 17:21:36.186940: I tensorflow/core/platform/cpu_feature_guard.cc:151] This TensorFlow binary is optimized with oneAPI Deep Neural
Network Library (oneDNN) to use the following CPU instructions in performance-critical operations: AVX AVX2
To enable them in other operations, rebuild TensorFlow with the appropriate compiler flags.
Epoch 1/10
560/570 [=====>.] - ETA: 0s - loss: 287.8470WARNING:tensorflow:Can save best model only with val_accuracy
available, skipping.
570/570 [=====] - 1s 1ms/step - loss: 283.3059
Epoch 2/10
527/570 [=====>...] - ETA: 0s - loss: 4.4827WARNING:tensorflow:Can save best model only with val_accuracy available,
skipping.
570/570 [=====] - 1s 1ms/step - loss: 5.2308
Epoch 3/10
565/570 [=====>.] - ETA: 0s - loss: 8.9520WARNING:tensorflow:Can save best model only with val_accuracy available,
skipping.
570/570 [=====] - 1s 1ms/step - loss: 8.8798
Epoch 4/10
568/570 [=====>.] - ETA: 0s - loss: 10.2503WARNING:tensorflow:Can save best model only with val_accuracy
available, skipping.
570/570 [=====] - 1s 931us/step - loss: 10.2216
Epoch 5/10
525/570 [=====>...] - ETA: 0s - loss: 10.9842WARNING:tensorflow:Can save best model only with val_accuracy available,
skipping.
570/570 [=====] - 1s 916us/step - loss: 10.1243
Epoch 6/10
518/570 [=====>...] - ETA: 0s - loss: 12.1647WARNING:tensorflow:Can save best model only with val_accuracy available,
skipping.
570/570 [=====] - 1s 918us/step - loss: 12.5059
Epoch 7/10
550/570 [=====>.] - ETA: 0s - loss: 14.5881WARNING:tensorflow:Can save best model only with val_accuracy available,
skipping.
570/570 [=====] - 1s 963us/step - loss: 14.7423
Epoch 8/10
566/570 [=====>.] - ETA: 0s - loss: 15.2827WARNING:tensorflow:Can save best model only with val_accuracy
available, skipping.
570/570 [=====] - 1s 933us/step - loss: 15.1862
Epoch 9/10
549/570 [=====>...] - ETA: 0s - loss: 9.3385WARNING:tensorflow:Can save best model only with val_accuracy available,
skipping.
570/570 [=====] - 1s 974us/step - loss: 9.0848
Epoch 10/10
549/570 [=====>...] - ETA: 0s - loss: 12.3400WARNING:tensorflow:Can save best model only with val_accuracy available,
skipping.
570/570 [=====] - 1s 1ms/step - loss: 11.8957
143/143 [=====] - 0s 743us/step - loss: 0.5875
EVALUATION: 0.5874761343002319
0.8024153660145087
0.9982444590739522
precision recall f1-score support
0 1.00 1.00 1.00 4549
1 0.50 0.38 0.43 8

accuracy 1.00 4557
macro avg 0.75 0.69 0.71 4557
weighted avg 1.00 1.00 1.00 4557

R2: -0.001758628269949325, MAE: 0.0017555409260478386, MSE: 0.0017555409260478386..
```